

CS3243: Introduction to Artificial Intelligence

Lecture Notes 3: Informed Search

1. Uniform-Cost Search (UCS) Algorithm

The Uniform-Cost Search algorithm improves on the Breadth-First Search (BFS) algorithm such that it is **optimal**.


- Instead of a Queue, a Priority Queue is used for the Frontier.
- $\hat{g}(u)$ keeps track of the minimum cost path of reaching node u discovered so far.

```
FindPathToGoal(u):  
    Frontier  $\leftarrow$  PriorityQueue(u)  
    Explored  $\leftarrow$  {u}  
     $\hat{g}[u] = 0$   
  
    while Frontier is not empty:  
        u  $\leftarrow$  Frontier.pop()  
        if GoalTest(u):  
            return path(u)  
        Explored.add(u)  
        for all children v of u:  
            if (v not in Explored):  
                if (v in Frontier):  
                     $\hat{g}[v] = \min(\hat{g}[v], \hat{g}[u] + c(u,v))$   
                else:  
                    Frontier.push(v)  
                     $\hat{g}[v] = \hat{g}[u] + c(u,v)$   
  
    return Failure
```

Note that UCS is different from Dijkstra's algorithm. In Dijkstra's algorithm, we initialize $\hat{g}(u)$ for all the nodes u .

- **Completeness:** If we assume that every edge of the graph has cost $\geq \epsilon$, then the algorithm will eventually reach the goal state.

However, this algorithm is not complete for graphs with edges of cost 0. The number of steps required to reach the goal state might potentially be infinite, and the algorithm may not be able to reach the goal. For example:



- **Optimality:** The algorithm is optimal. When we perform the Goal Test on a node u , we have already found the optimal path to u .

Proof. We will prove that when we pop a node u from the frontier, we have found the optimal path to u .

We make the following definitions:

- $g(u)$ is the actual minimum path cost from the start node to node u .
- $\hat{g}(u)$ is the minimum path cost from the start node to node u discovered so far.
- $\hat{g}_{pop}(u)$ is the value of $\hat{g}(u)$ when we pop u from the frontier.

Let us consider the optimal path $\pi: S_0, S_1, S_2, \dots, S_k, u$.

We will perform induction on the path.

Base Case:

When we pop node S_0 from the frontier,

$$\hat{g}_{pop}(S_0) = g(S_0) = 0 \text{ as } S_0 \text{ is the start node.}$$

$$\hat{g}_{pop}(S_1) = g(S_1) \text{ as ensured by the Priority Queue.}$$

$$\hat{g}(S_1) = \min(\hat{g}(S_1), \hat{g}(S_0) + c(S_0, S_1)) \text{ as ensured by the algorithm.}$$

Observe that $g(S_0) \leq g(S_1) \leq g(S_2) \leq \dots \leq g(u)$.

Hence, at all times,

$$\hat{g}(u) \geq g(u) \geq g(S_k) \text{ --- (1)}$$

Inductive Hypothesis:

Let us assume that for all nodes $\{S_0, \dots, S_k\}$, we find that $\hat{g}_{pop}(S_i) = g(S_i)$.

Then, when we pop S_k , $\hat{g}(u) \leq \hat{g}_{pop}(S_k) + c(S_k, u) = g(S_k) + c(S_k, u) = g(u)$

Hence,

$$\hat{g}(u) \leq g(u) \text{ --- (2)}$$

From (1) and (2), we can conclude that

$$\hat{g}_{pop}(u) = g(u)$$

- **Time:** $O(b^{1+d})$. Assuming that the optimal cost to reach the goal node is C^* , and each edge has cost $\geq \epsilon$. Then $\frac{C^*}{\epsilon} = d$.
- **Space:** $O(b^{1+d})$ which is still expensive.

2. Depth-First Search (DFS) Algorithm

The Depth-First Search (DFS) algorithm is just BFS but with the frontier changed to a last-in-first-out stack. This reduces the space complexity to $O(bd)$. However, the tradeoff is that the algorithm is not complete.

3. A* Algorithm

BFS, DFS and UCS are all *uninformed* search algorithms where there is no information on how much it might actually cost to reach a goal. In contrast, the A* algorithm is an *informed* search algorithm.

This algorithm is similar to UCS, except that it makes use of the function $\hat{f}(u)$ instead of $\hat{g}(u)$, where:

- $\hat{g}(u)$ is the minimum path cost from the start node to node u discovered so far.
- $\hat{h}(u)$, or the **heuristic**, is an estimate of the path cost from node u to the goal node.
- $\hat{f}(u) = \hat{g}(u) + \hat{h}(u)$

```
FindPathToGoal(u):  
    // PriorityQueue to be implemented with  $\hat{f}$   
    Frontier  $\leftarrow$  PriorityQueue(u)  
    Explored  $\leftarrow$  {u}  
     $\hat{g}[u] = 0$   
  
    while Frontier is not empty:  
        u  $\leftarrow$  Frontier.pop()  
        if GoalTest(u):  
            return path(u)  
        Explored.add(u)  
        for all children v of u:  
            if (v not in Explored):  
                if (v in Frontier):  
                     $\hat{g}[v] = \min(\hat{g}[v], \hat{g}[u] + c(u,v))$   
                     $\hat{f}(v) = \hat{g}(v) + \hat{h}(v)$   
                else:  
                    Frontier.push(v)  
                     $\hat{g}[v] = \hat{g}[u] + c(u,v)$   
                     $\hat{f}(v) = \hat{g}(v) + \hat{h}(v)$   
  
    return Failure
```

What is the property of $h(u)$ that makes sure that this algorithm is optimal?

Recall: UCS was optimal because on the optimal path $\pi: S_0, S_1, S_2, \dots, S_k, u$, we have:

1. $\hat{g}_{pop}(S_0) \leq \hat{g}_{pop}(S_1) \leq \dots \leq \hat{g}_{pop}(S_k) \leq \hat{g}_{pop}(u)$
2. $\hat{g}_{pop}(S_i) = g(S_i)$

Similarly, the A* algorithm is optimal if we have the following properties:

P1: $\hat{f}_{pop}(S_0) \leq \hat{f}_{pop}(S_1) \leq \dots \leq \hat{f}_{pop}(S_k) \leq \hat{f}_{pop}(u)$

P2: $\hat{f}_{pop}(S_i) = f(S_i)$

Properties of the Heuristic, $\hat{h}(u)$

In order to ensure that **P1** and **P2** holds (such that the A* algorithm is optimal), the heuristic must fulfill the following properties:

- **Consistency:** $h(S_i) \leq c(S_i, S_{i+1}) + h(S_{i+1})$, also known as the triangle inequality.

Derivation. If $f(S_0) \leq f(S_1) \leq \dots \leq f(S_k) \leq f(u)$ is true, and **P2** holds, then **P1** would hold. Hence, we want to make sure that $f(S_i) \leq f(S_{i+1})$.

Expanding and manipulating the above inequality,

$$\begin{aligned} g(S_i) + h(S_i) &\leq g(S_{i+1}) + h(S_{i+1}) \\ h(S_i) &\leq g(S_{i+1}) - g(S_i) + h(S_{i+1}) \end{aligned}$$

Therefore, we get $h(S_i) \leq c(S_i, S_{i+1}) + h(S_{i+1})$.

- **Admissibility:** $h(S_i) \leq OPT(S_i)$, where $OPT(S_i)$ refers to the optimal path cost from node.

Derivation: Expanding from the property of consistency,

$$\begin{aligned} h(S_i) &\leq c(S_i, S_{i+1}) + h(S_{i+1}) \\ h(S_i) &\leq c(S_i, S_{i+1}) + c(S_{i+1}, S_{i+2}) + h(S_{i+2}) \\ &\quad \vdots \\ h(S_i) &\leq c(S_i, S_{i+1}) + c(S_{i+1}, S_{i+2}) + \dots + c(S_k, u) + h(u) \end{aligned}$$

Since $h(u) = 0$ as u is the goal node,

$$h(S_i) \leq c(S_i, S_{i+1}) + c(S_{i+1}, S_{i+2}) + \dots + c(S_k, u) = OPT(S_i)$$

Therefore, we get $h(S_i) \leq OPT(S_i)$.